

Implementation of the Apriori Algorithm

Raiyan Abdul Baten
University of Rochester
Rochester, NY
rbaten@ur.rochester.edu

ABSTRACT

This document explains my approach for mining frequent itemsets from the adult dataset. In particular, the report sheds light on data preparation, python implementation and result analysis of the Apriori algorithm.

KEYWORDS

Apriori, Frequent itemset mining

ACM Reference Format:

Raiyan Abdul Baten. 2018. Implementation of the Apriori Algorithm. In *Proceedings of ACM Conference (Conference'17)*. ACM, New York, NY, USA, 1 page. <https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

1 INTRODUCTION

Apriori is a popular algorithm for mining frequent itemsets from transaction databases. In this project, I have implemented the algorithm as specified in Chapter 6 of Han et al.'s book [1].

2 DATASET

I have used the adult dataset¹ in this project.

2.1 Dataset Summary

The dataset was originally intended for classification tasks, and therefore comes pre-split into training and test sets, in files named 'adult.data' and 'adult.test' respectively. The training set contains $N_{\text{training}} = 32561$ data tuples, and the test set contains $N_{\text{test}} = 16281$ tuples. In this project, I have analyzed the two sets separately to obtain the results.

Both of the dataset files contain the same 14 attributes and 1 class label: age, workclass, fnlwgt, education, education-num, marital-status, occupation, relationship, race, sex, capital-gain, capital-loss, hours-per-week, native-country, and the class label of yearly income. Since this is a frequent itemset mining task rather than classification, I have treated all of the attributes and class labels as 'items', making all 15 of the information entries conceptually equivalent.

¹<http://archive.ics.uci.edu/ml/datasets/Adult>

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

Conference'17, July 2017, Washington, DC, USA

© 2018 Copyright held by the owner/author(s).

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM.

<https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

2.2 Pre-processing

A quick exploration reveals that the attributes 'education' and 'education-num' have one-to-one mapping. For example, 'Bachelors' in 'education' is coded with '13' in 'education-num', 'Masters' is coded with '14', so on and so forth. Hence, the attribute 'education-num' is just a proxy for 'education', and therefore I removed the 'education-num' attribute altogether from my analyses.

The attributes 'workclass', 'occupation' and 'native-country' have blank entries in them, labeled with '?' marks. I decided to treat the '?' entries as categorical values of the respective attributes, since all of the entries are eventually treated as 'items' anyway. To differentiate between the '?' entries of one attribute from those of others, I concatenated all of the entries of these three attributes with their respective attribute names. Therefore, the workclass 'Private' became 'Private_workclass', the entry '?' became '?_workclass', so on and so forth for occupation and native-country as well.

The attributes 'age', 'fnlwgt', 'capital-gain', 'capital-loss' and 'hours-per-week' are numeric. I split each of them into 10 equal-width bins, and took the bin boundary information as categorical values. I further concatenated these categorical names with the attributes names for clarity. Therefore, the age '39' became 'age_(38.9,46.2]', holding the attribute and binning information.

Finally, I converted the dataset into a dictionary, with keys corresponding to 'transaction ID's, and values corresponding to the lists of 'items' (holding attribute values). Therefore, every tuple in the original dataset became a (processed) list of items. This dictionary was fed into the code as the input database.

3 IMPLEMENTATION

I implemented the algorithm presented in Figure 6.4 of Han et al.'s textbook [1]. I chose a support of 23%.

Data structures: Each 'item' is represented with a string. I have used the list data structure in python to store each itemset. A list of lists is used to store the frequent k -itemsets $\forall k \in \{1, 2, 3, \dots\}$. Then, all of these k -itemsets are further stored in a list of lists of lists, named L . Therefore, $L[0]$ holds frequent 1-itemsets, $L[1]$ holds frequent 2-itemsets, so on and so forth. I have also used a list of lists of dictionaries, C , to hold the support information along with the itemsets. Therefore, $C[0]$ is a list of dictionaries, where each of the dictionaries hold the itemset members and supports of the frequent 1-itemsets.

4 RESULTS

The 'adult.data' and 'adult.test' datasets gave 581 and 587 frequent itemsets respectively, with a maximum $k = 7$.

REFERENCES

- [1] Jiawei Han, Jian Pei, and Micheline Kamber. 2011. *Data mining: concepts and techniques*. Elsevier.